

Assignment 2

Bayesian Inference FTN0548

Usama Zafar, Li Ju

March 4, 2024

1 Problem 4.1

(a) With the given $p(x|\theta)$ in the table and $\pi(\theta)$, the joint distribution and marginal distributions of x and θ are computed as follows and shown in Table 1.

$$p(x, \theta) = p(x|\theta)\pi(\theta)$$

$$p(x) = \sum_{\theta \in \{0,1,2\}} p(x, \theta)$$

	$x = -1$	$x = 0$	$x = 1$	$\pi(\theta)$
$\theta = 0$	$0.5 - 0.5p$	$0.3 - 0.3p$	$0.2 - 0.2p$	$1 - p$
$\theta = 1$	0	$0.5p$	$0.5p$	p
$p(x)$	$0.5 - 0.5p$	$0.3 + 0.2p$	$0.2 + 0.3p$	

Table 1: Joint and marginal distributions of x and θ .

The posterior distribution of θ given x is computed as follows and shown in Table 2.

$$\pi(\theta|x) = \frac{p(x, \theta)}{p(x)}$$

	$x = -1$	$x = 0$	$x = 1$
$\theta = 0$	1	$\frac{3-3p}{3+2p}$	$\frac{2-2p}{2+3p}$
$\theta = 1$	0	$\frac{5p}{3+2p}$	$\frac{5p}{2+3p}$

Table 2: Posterior distribution of θ given x $\Pr(\theta|x)$.

(b) The Bayes estimator is defined as

$$\hat{\theta}^\pi := \arg \min_{\hat{\theta}} \mathbb{E}_{\pi(\theta|x)} L(\theta, \hat{\theta})$$

$$= \arg \min_{\hat{\theta}} \sum_{\theta_0 \in \{0,1\}} L(\theta_0, \hat{\theta}) \Pr(\theta = \theta_0|x)$$

The expected posterior loss is computed as follows:

$$\mathbb{E}_{\pi(\theta|x)} L(\theta, \hat{\theta}) = \begin{cases} L(0, \hat{\theta}) & x = -1 \\ L(0, \hat{\theta})^{\frac{3-3p}{3+2p}} + L(1, \hat{\theta})^{\frac{5p}{3+2p}} & x = 0 \\ L(0, \hat{\theta})^{\frac{2-2p}{2+3p}} + L(1, \hat{\theta})^{\frac{5p}{2+3p}} & x = 1 \end{cases}$$

Hence, with the 0 – 1 loss function L , the Bayes estimator for θ is given in Table 3.

	$p > \frac{3}{8}$	$\frac{2}{7} \geq p > \frac{2}{7}$	$\frac{2}{7} \geq p$
$x = -1$	0	0	0
$x = 0$	1	0	0
$x = 1$	1	1	0

Table 3: The Bayes estimator for θ .

(c) The frequentist risk is defined by

$$\begin{aligned} R(\theta_0, \hat{\theta}) &:= \mathbb{E}_{p(x|\theta_0)} L(\theta_0, \hat{\theta}) \\ &= \sum_{x \in \{-1, 0, 1\}} L(\theta_0, \hat{\theta}) \Pr(X = x | \theta_0) \end{aligned}$$

Hence, the frequentist risks of the Bayes estimator of θ are given in Table 4.

	$p > \frac{3}{8}$	$\frac{2}{7} \geq p > \frac{2}{7}$	$\frac{2}{7} \geq p$
$\theta_0 = 0$	0.5	0.2	0
$\theta_0 = 1$	0	0.5	1

Table 4: The frequentist risk for the Bayes estimator $\hat{\theta}^\pi$.

(d) The Bayes risk is defined by

$$r(\pi, \delta) := \mathbb{E}^\pi R(\theta, \delta) = \sum_{\theta_0 \in \{0, 1\}} R(\theta_0, \delta) \pi(\theta_0)$$

Hence, the Bayes risk of the Bayes estimator for θ is given by

$$r(\pi, \hat{\theta}^\pi) = \begin{cases} 0.5 - 0.5p & p > \frac{3}{8} \\ 0.2 + 0.3p & \frac{2}{7} \geq p > \frac{2}{7} \\ p & \frac{2}{7} \geq p \end{cases} \quad (1)$$

(e) The least favourable prior is defined by

$$\pi_0 := \arg \max_{\pi} r(\pi, \hat{\theta}^\pi)$$

Given that $r(\pi, \hat{\theta}^{\pi(p)})$ is a piecewise linear function w.r.t p as shown in Equation 1, the supremum of the $r(\pi, \hat{\theta}^{p_i})$ is given when $p = \frac{3}{8}$. Hence, the least favourable prior is $\pi_0 = \Pr(\theta = 1) = \frac{3}{8}$.

2 Problem 4.2

(a) Recognizing that the prior is conjugate to the binomial likelihood function, with observation x , we have the posterior distribution for θ as

$$\pi(\theta|x) \sim \text{Beta}(\alpha + x, \beta + n - x)$$

Then the expected posterior loss is given by

$$\begin{aligned} \mathbb{E}_{p(\theta|x)} L(\theta, d) &= \int L(\theta, d) \theta^{\alpha+x-1} (1-\theta)^{\beta+n-x+1} d\theta \cdot \frac{1}{\mathbf{B}(\alpha+x, \beta+n-x)} \\ &= k_1 \int_0^d (d-\theta) \theta^{\alpha+x-1} (1-\theta)^{\beta+n-x+1} d\theta \cdot \frac{1}{\mathbf{B}(\alpha+x, \beta+n-x)} \\ &\quad + k_2 \int_d^1 (\theta-d) \theta^{\alpha+x-1} (1-\theta)^{\beta+n-x+1} d\theta \cdot \frac{1}{\mathbf{B}(\alpha+x, \beta+n-x)} \\ &= k_1 \left(d \text{CDF}_{\alpha+x, \beta+n-x}(d) - \text{CDF}_{\alpha+x+1, \beta+n-x}(d) \frac{\mathbf{B}(\alpha+x+1, \beta+n-x)}{\mathbf{B}(\alpha+x, \beta+n-x)} \right) \\ &\quad + k_2 \left((1 - \text{CDF}_{\alpha+x+1, \beta+n-x}(d)) \frac{\mathbf{B}(\alpha+x+1, \beta+n-x)}{\mathbf{B}(\alpha+x, \beta+n-x)} - d(1 - \text{CDF}_{\alpha+x, \beta+n-x}(d)) \right) \\ &= k_1 \left(d \text{CDF}_{\alpha+x, \beta+n-x}(d) - \text{CDF}_{\alpha+x+1, \beta+n-x}(d) \frac{\alpha+x}{\alpha+\beta+n} \right) \\ &\quad + k_2 \left((1 - \text{CDF}_{\alpha+x+1, \beta+n-x}(d)) \frac{\alpha+x}{\alpha+\beta+n} - d(1 - \text{CDF}_{\alpha+x, \beta+n-x}(d)) \right) \end{aligned}$$

where $\mathbf{B}(a, b) = \frac{\Gamma(a)\Gamma(b)}{\Gamma(a+b)}$, and $\text{CDF}_{a,b}$ denotes the cumulative distribution function of distribution $\text{Beta}(a, b)$.

The derivative of $\mathbb{E}_{p(\theta|x)} L(\theta, d)$ w.r.t d is given by

$$\begin{aligned} \frac{\partial \mathbb{E}_{p(\theta|x)} L(\theta, d)}{\partial d} &= k_1 \left(\text{CDF}_{\alpha+x, \beta+n-x}(d) + d \text{PDF}_{\alpha+x, \beta+n-x}(d) - \frac{\alpha+x}{\alpha+\beta+n} \text{PDF}_{\alpha+x+1, \beta+n-x}(d) \right) \\ &\quad + k_2 \left(-\frac{\alpha+x}{\alpha+\beta+n} \text{PDF}_{\alpha+x+1, \beta+n-x}(d) - 1 + \text{CDF}_{\alpha+x, \beta+n-x}(d) + d \text{PDF}_{\alpha+x, \beta+n-x}(d) \right) \\ &= -k_2 + (k_1 + k_2) \text{CDF}(d) \end{aligned}$$

where $\text{PDF}_{a,b}$ denotes the probability density function of $\text{Beta}(a, b)$.

The Bayes estimator is given when $\frac{\partial \mathbb{E}_{p(\theta|x)} L(\theta, d)}{\partial d}(\hat{\theta}^\pi) = 0$. Hence, we have

$$\text{CDF}_{\alpha+x, \beta+n-x}(\hat{\theta}^\pi) = \frac{k_2}{k_1 + k_2}$$

Hence, the Bayes estimator for θ is given by the $\frac{k_2}{k_1+k_2}$ quantile of the posterior distribution $\text{Beta}(\alpha+x, \beta+n-x)$.

(b) Let $k_1 = k_2 = k$. The Bayes estimator satisfies

$$\text{CDF}_{\alpha+x, \beta+n-x}(\hat{\theta}^\pi) = \frac{1}{2}$$

Recall that the median of $\text{Beta}(a, b)$ is approximately given by $\frac{a-\frac{1}{3}}{a+b-\frac{2}{3}}$, the Bayes estimator is given by

$$\hat{\theta}^\pi(x) = \frac{\alpha+x-\frac{1}{3}}{\alpha+\beta+n-\frac{2}{3}}$$

(c) The use of asymmetric loss is advantageous when the repercussions of overestimation differ significantly from those of underestimation. By allocating larger coefficients to the components that could lead to more serious consequences, the frequency of such adverse outcomes may be diminished. An real-life example is medical diagnosis, in which the use of an asymmetric loss function is crucial because the costs of false negatives and false positives are highly imbalanced.

(d) The risk for estimator $\hat{\theta}$ given θ_0 is defined as

$$\begin{aligned} R(\theta_0, \hat{\theta}) &:= \mathbb{E}_{p(X|\theta_0)} L(\theta_0, \hat{\theta}) \\ &= \sum_{x=0}^n L(\theta_0, \hat{\theta}) \binom{n}{x} \theta_0^x (1 - \theta_0)^{n-x} \end{aligned}$$

tobecontinued

3 Problem 5.1

(a) Given the likelihood function $p(X|\theta) = \text{Exp}(\theta)$ and $\pi(\theta) = \Gamma(\alpha, \beta)$, we have the posterior distribution of θ given an i.i.d. sample $\mathbf{x}_{(n)} = (x_1, \dots, x_n)$

$$\begin{aligned} \pi(\theta|\mathbf{x}_{(n)}) &\propto p(\mathbf{x}_{(n)}|\theta)\pi(\theta) \\ &= \prod \text{Exp}(\theta)\Gamma(\alpha, \beta) \\ &= \theta^n \exp\{-\theta \sum_{i=1}^n x_i\} \theta^{\alpha-1} \exp\{-\beta\theta\} \\ &= \theta^{n+\alpha-1} \exp\{-\theta(\beta + \sum_{i=1}^n x_i)\} \end{aligned}$$

Hence, the posterior distribution $\pi(\theta|\mathbf{x}_{(n)}) = \Gamma(n + \alpha, \beta + \sum_{i=1}^n x_i)$.

(b) With L_2 loss, the Bayes estimator for θ is given by

$$\begin{aligned} \delta(\mathbf{x}_{(n)}) &= \arg \min_{\hat{\theta}} \mathbb{E}_{\pi(\theta|\mathbf{x}_{(n)})} L_2(\theta, \hat{\theta}) \\ &= \arg \min_{\hat{\theta}} \int (\theta - \hat{\theta})^2 \theta^{n+\alpha-1} \exp\{-\theta(\beta + \sum_{i=1}^n x_i)\} d\theta \cdot \frac{(\beta + \sum_{i=1}^n x_i)^{\alpha+n}}{\Gamma(n + \alpha)} \\ &= \arg \min_{\hat{\theta}} \hat{\theta}^2 - 2\hat{\theta} \int \theta^{n+\alpha} \exp\{-\theta(\beta + \sum_{i=1}^n x_i)\} d\theta \cdot \frac{(\beta + \sum_{i=1}^n x_i)^{\alpha+n}}{\Gamma(n + \alpha)} + \text{const.} \\ &= \arg \min_{\hat{\theta}} \hat{\theta}^2 - 2\hat{\theta} \frac{\Gamma(n + \alpha + 1)}{(\beta + \sum_{i=1}^n x_i)^{\alpha+n+1}} \frac{(\beta + \sum_{i=1}^n x_i)^{\alpha+n}}{\Gamma(n + \alpha)} + \text{const.} \\ &= \arg \min_{\hat{\theta}} \hat{\theta}^2 - 2\hat{\theta} \frac{n + \alpha}{\beta + \sum_{i=1}^n x_i} + \text{const.} \\ \delta(\mathbf{x}_{(n)}) &= \frac{n + \alpha}{\beta + \sum_{i=1}^n x_i} \end{aligned}$$

(c) tobecontinued

4 Problem 5.2

(a) Assume $p(\boldsymbol{\theta})$ in exponential family in the form of $p(\boldsymbol{\theta}) = h(x)g(\boldsymbol{\theta})\exp(\boldsymbol{\eta}(\boldsymbol{\theta})^\top \mathbf{T}(x))$. The Kullback-Leibler divergence between $p(\boldsymbol{\theta}_0)$ and $p(\boldsymbol{\theta})$ is given as follows:

$$\begin{aligned}\text{KL}(p(\boldsymbol{\theta}_0)||p(\boldsymbol{\theta})) &= \int p_x(\boldsymbol{\theta}_0) \ln \frac{p_x(\boldsymbol{\theta}_0)}{p_x(\boldsymbol{\theta})} dx \\ &= \int p(\boldsymbol{\theta}_0) \ln \frac{h(x)g(\boldsymbol{\theta}_0)\exp(\boldsymbol{\eta}(\boldsymbol{\theta}_0)^\top \mathbf{T}(x))}{h(x)g(\boldsymbol{\theta})\exp(\boldsymbol{\eta}(\boldsymbol{\theta})^\top \mathbf{T}(x))} dx \\ &= \ln \frac{g(\boldsymbol{\theta}_0)}{g(\boldsymbol{\theta})} + \int p(\boldsymbol{\theta}_0) (\boldsymbol{\eta}(\boldsymbol{\theta}_0) - \boldsymbol{\eta}(\boldsymbol{\theta}))^\top \mathbf{T}(x) dx \\ &= \ln \frac{g(\boldsymbol{\theta}_0)}{g(\boldsymbol{\theta})} + (\boldsymbol{\eta}(\boldsymbol{\theta}_0) - \boldsymbol{\eta}(\boldsymbol{\theta}))^\top \mathbb{E}_{p(\boldsymbol{\theta}_0)} [\mathbf{T}(x)]\end{aligned}$$

(b) For Gamma distributions, first we reform $\Gamma(\alpha, \beta)$ in exponential form:

$$\Gamma_x(\alpha, \beta) = h(x)g(\alpha, \beta)\exp(\boldsymbol{\eta}(\alpha, \beta)^\top \mathbf{T}(x))$$

where $h(x) = 1$, $g(\boldsymbol{\theta}) = \frac{\beta^\alpha}{\Gamma(\alpha)}$, $\boldsymbol{\eta}(\alpha, \beta) = (\alpha - 1, -\beta)^\top$ and $\mathbf{T}(x) = (\ln x, x)^\top$.

Then the Kullback-Leibler divergence between $\Gamma(\alpha_0, \beta_0)$ and $\Gamma(\alpha, \beta)$ is given by

$$\begin{aligned}\text{KL}(\Gamma(\alpha_0, \beta_0)||\Gamma(\alpha, \beta)) &= \ln \frac{\beta_0^{\alpha_0} \Gamma(\alpha)}{\beta^\alpha \Gamma(\alpha_0)} + (\alpha_0 - \alpha) \mathbb{E}_{\Gamma(\alpha_0, \beta_0)}[\ln x] - (\beta_0 - \beta) \mathbb{E}_{\Gamma(\alpha_0, \beta_0)}[x] \\ &= \alpha_0 \ln \beta_0 - \alpha \ln \beta + \ln \Gamma(\alpha) - \ln \Gamma(\alpha_0) + (\alpha_0 - \alpha)(\psi(\alpha_0) - \ln \beta_0) - (\beta_0 - \beta) \frac{\alpha_0}{\beta_0} \\ &= \alpha(\ln \beta_0 - \ln \beta) + \ln \Gamma(\alpha) - \ln \Gamma(\alpha_0) + (\alpha_0 - \alpha)\psi(\alpha_0) + \frac{\beta - \beta_0}{\beta_0} \alpha_0\end{aligned}$$

where ψ is the digamma function $\psi(z) = \frac{d\Gamma(z)}{dz}$.

5 Problem 5.4

(a) tobecontinued

6 Problem 6.1

(a)

7 Problem 6.2

(a)

8 Problem 6.3

(a)